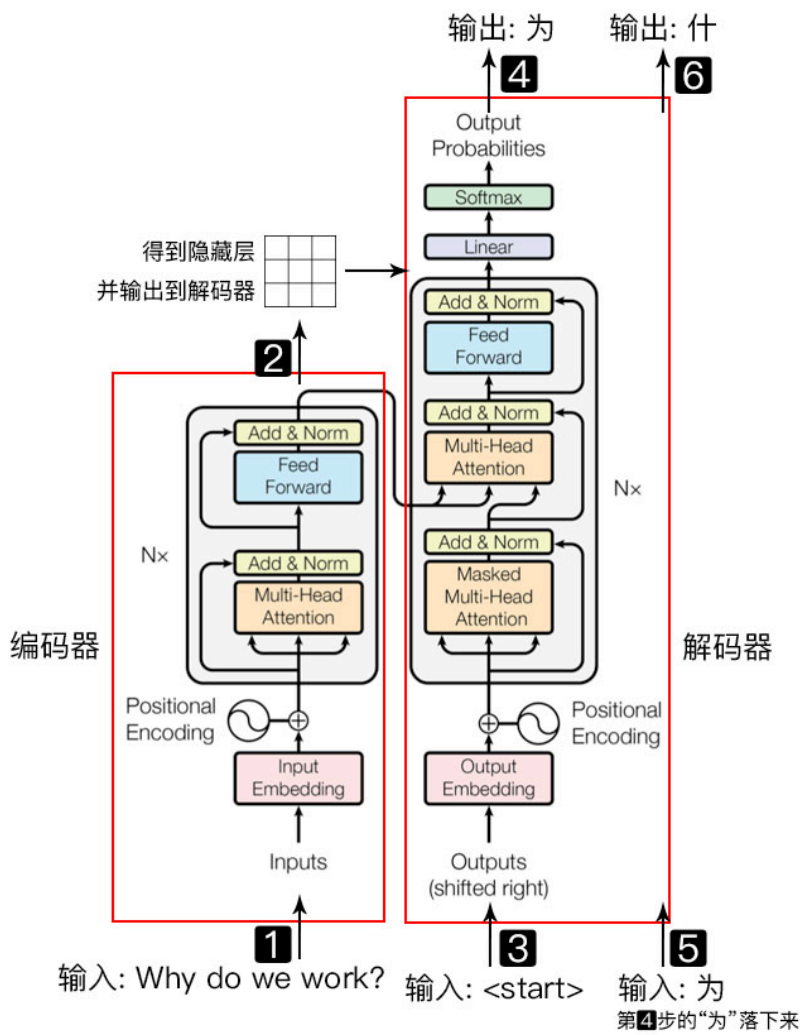


整体架构

过程

编码器负责把自然语言序列映射成为隐藏层， 解码器把隐藏层再映射为自然语言序列。

- 1. 输入自然语言序列到编码器：Why do we work?（为什么要工作）；
- 2. 通过编码器得到隐藏层，再输入到解码器；
- 3. 输入<start>（起始）符号到解码器；
- 4. 得到第一个字"为"；
- 5. 将得到的第一个字"为"落下来再输入到解码器；
- 6. 得到第二个字"什"；
- 7. 将得到的第二字再落下来，直到解码器输出<end>（终止符），即序列生成完成。



Transformer

```
1 class Transformer(nn.Module):
2     # model_dim: Embedding Size, ffn_dim: FeedForward dimension
```

```

3      # n_layers: number of Encoder of Decoder Layer, n_heads: number of
      heads in Multi-Head Attention
4      def __init__(self, src_vocab_size, src_max_len, tgt_vocab_size,
      tgt_max_len,
5                      num_layers=6, num_heads=8, model_dim=512, ffn_dim=2048,
      dropout=0.0):
6          super(Transformer, self).__init__()
7
8          self.encoder = Encoder(src_vocab_size, src_max_len, num_layers,
      model_dim, num_heads, ffn_dim, dropout)
9          self.decoder = Decoder(tgt_vocab_size, tgt_max_len, num_layers,
      model_dim, num_heads, ffn_dim, dropout)
10         self.linear = nn.Linear(model_dim, tgt_vocab_size, bias=False)
11         self.softmax = nn.Softmax(dim=2)
12
13         def forward(self, src_seq, tgt_seq):
14             src_mask = get_pad_mask(src_seq).int()
15             tgt_mask = torch.gt((get_pad_mask(tgt_seq).int() +
      get_sequence_mask(tgt_seq).int()), 0).int()
16
17             enc_output, enc_self_attn = self.encoder(src_seq, src_mask)
18             dec_output, dec_self_attn, dec_enc_attn = self.decoder(tgt_seq,
      tgt_mask, enc_output, src_mask)
19             output = self.softmax(self.linear(dec_output))
20
21             return output, enc_self_attn, dec_self_attn, dec_enc_attn

```

Mask

Transformer 模型里面涉及两种mask，分别是 padding mask 和 sequence mask。

padding mask 由于对输入设置了最大的输入长度（max sequence length），因此对短于最大长度的输入序列会在其后面填充0。但是这些填充的位置并没无意义，attention 机制不应该把注意力放在这些位置上，因此可以把 $mask = 1$ 对应的输入位置加上一个非常大的负数（或负无穷），这样经过 softmax 这些位置的概率就会接近0。

```

1  def get_pad_mask(seq):
2      mask = (seq == 0).unsqueeze(-2)
3      return mask

```

sequence mask sequence mask 是为了使 decoder 不能看到未来的信息。对于一个序列，在 t 时刻解码输出只能依赖于 t 时刻之前的输出，而不能依赖 t 之后的输出。可以构造一个矩阵，上三角的值全为1，下三角的值全为0，对角线也是0，值为1的位置即是加上负无穷的位置。

```

1 def get_sequence_mask(seq):
2     batch_size, seq_len = seq.size()
3     mask = torch.triu(torch.ones((seq_len, seq_len), dtype=torch.uint8),
4     diagonal=1)
5     mask = mask.unsqueeze(0).expand(batch_size, -1, -1)
6     return mask

```

Positional Encoding

位置编码用于提供语言的位置顺序信息，论文使用了 sin 和 cos 函数的变换：

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$

位置嵌入的维度为 $(\text{max sequence length}, \text{embedding size})$ ， pos 指的是句中字的位置，取值范围是 $[0, \text{max sequence length})$ ， i 指的是词向量的维度，取值范围是 $[0, \text{embedding size})$ 。


```

1 class PositionalEncoding(nn.Module):
2     def __init__(self, max_seq_length, embed_size):
3         super(PositionalEncoding, self).__init__()
4
5         pe = np.array([[pos / np.power(10000, 2 * (i // 2) / embed_size))
6         for i in range(d_model)] for pos in
7         range(max_seq_length)])
8
9         pe[:, 0::2] = np.sin(pe[:, 0::2])
10        pe[:, 1::2] = np.cos(pe[:, 1::2])
11        pe = torch.FloatTensor(pe).unsqueeze(0) # (1, max_seq_length,
12        d_model)
13        self.register_buffer("pe", pe)
14
15    def forward(self, x):
16        # x: input--(batch size, max sequence length)
17        # (batch size, max seq length, embedding size)
18        return self.pe[:, :x.size(1), :].clone().detach()

```

Mutil-Head Attention

Scaled Dot-Product Attention



$X_{embedding} = EmbeddingLookup(X) + PositionalEncoding(X)$, 其维度为 $(batch\ size, max\ sequence\ length, embedding\ size)$.

$$Attention(Q, K, V) = softmax(\frac{QK^T}{\sqrt{d_k}})V$$

- 在 encoder 的 self-attention 中，Q、K、V 都来自上一层 encoder 的输出。对于第一层 encoder，其输入为 $X_{embedding}$ 。
- 在 decoder 的 self-attention 中，Q、K、V 都来自上一层 decoder 的输出。对于第一层 decoder，其输入为 $X_{embedding}$ ，但是对于 decoder，不希望它能获得下一个时间步的信息，因此需要进行 sequence masking。
- 在 encoder-decoder attention 中，Q 来自 decoder 的上一层输出，K、V 来自 encoder 的输出。

```

1  class ScaledDotProductAttention(nn.Module):
2      def __init__(self, dropout=0.2):
3          super(ScaledDotProductAttention, self).__init__()
4
5          self.softmax = nn.Softmax(dim=2)
6          self.dropout = nn.Dropout(dropout)
7
8      def forward(self, q, k, v, scale=None, mask=None):
9          # q,k,v (batch size, max seq length, embedding size)
10         # attention (batch size, max seq length, max seq length)
11         attention = torch.bmm(q, k.transpose(1, 2))
12         if scale:
13             attention = attention * scale
14         if mask is not None:
15             attention = attention.masked_fill_(mask.bool(), -np.inf)


```

```

16         attention = self.dropout(self.softmax(attention))
17         # context (batch size, max seq length, embedding size)
18         context = torch.bmm(attention, v)
19
20         return context, attention

```

Mutil-Head Attention



$$Q = \text{Linear}(X_{\text{embedding}}) = X_{\text{embedding}} W_Q$$

$$K = \text{Linear}(X_{\text{embedding}}) = X_{\text{embedding}} W_K$$

$$V = \text{Linear}(X_{\text{embedding}}) = X_{\text{embedding}} W_V$$

$$\text{MutilHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_n) W_O$$

Q, K, V (batch size, max sequence length, embedding size), W_Q, W_K, W_V, W_O (embedding size, embedding size).

将 Q, K, V 通过一个线性映射之后把 *embedding size* 分成 h 份，对每一份进行 Scaled Dot-Product Attention。然后，把各个部分的结果合并起来再经过线性映射，得到最终的输出。

```

1 class MultiHeadAttention(nn.Module):
2     def __init__(self, model_dim=512, num_heads=8, dropout=0.2):
3         super(MultiHeadAttention, self).__init__()
4


```

```

5         self.dim_per_head = model_dim // num_heads
6         self.num_heads = num_heads
7         self.linear_q = nn.Linear(model_dim, self.dim_per_head *
num_heads)
8         self.linear_k = nn.Linear(model_dim, self.dim_per_head *
num_heads)
9         self.linear_v = nn.Linear(model_dim, self.dim_per_head *
num_heads)
10        self.attention = ScaledDotProductAttention(dropout)
11        self.linear = nn.Linear(model_dim, model_dim)
12        self.dropout = nn.Dropout(dropout)
13        self.layer_norm = nn.LayerNorm(model_dim)
14
15    def forward(self, q, k, v, mask=None):
16        residual = q
17
18        dim_per_head = self.dim_per_head
19        num_heads = self.num_heads
20        batch_size = q.size(0)
21
22        # linear projection
23        q = self.linear_q(q)
24        k = self.linear_k(k)
25        v = self.linear_v(v)
26
27        # split by num_heads
28        q = q.view(batch_size * num_heads, -1, dim_per_head)
29        k = k.view(batch_size * num_heads, -1, dim_per_head)
30        v = v.view(batch_size * num_heads, -1, dim_per_head)
31
32        if mask is not None:
33            mask = mask.repeat(num_heads, 1, 1)
34
35        scale = (q.size(-1) // num_heads) ** -0.5
36        context, attention = self.attention(q, k, v, scale, mask)
37
38        # concat heads
39        context = context.view(batch_size, -1, dim_per_head * num_heads)
40
41        output = self.dropout(self.linear(context))
42        output = self.layer_norm(residual + output)
43
44        return output, attention

```

Attention 的含义



$$\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

先计算 Q 、 K 的点积，两个向量越相似，其点积越大。比如第一个字与第一个字的点积，即 $c1$ 行与 $c1$ 列相乘得到的 $c1c1$ ，代表了第一个字的注意力机制，注意力矩阵的第一行指的是第一个字与所有6个字的相似度。

然后再点乘 V 的列，从而使得每个字向量都含有当前句子所有字向量的全部信息。

Position-wise FeedForward Network

这是一个全连接网络，包含两个线性变换和一个非线性函数（ReLU）。

$$FFN(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

输入 x ($batch\ size, max\ sequence\ length, model\ dim$)

W_1 ($model\ dim, ffn\ dim$)

W_2 ($ffn\ dim, model\ dim$)

输出 ($batch\ size, max\ sequence\ length, model\ dim$)

```
1 class PositionalWiseFeedForward(nn.Module):
2     def __init__(self, model_dim=512, ffn_dim=2048, dropout=0.2):
3         super(PositionalWiseFeedForward, self).__init__()
4
5         self.linear1 = nn.Linear(model_dim, ffn_dim)
```

```

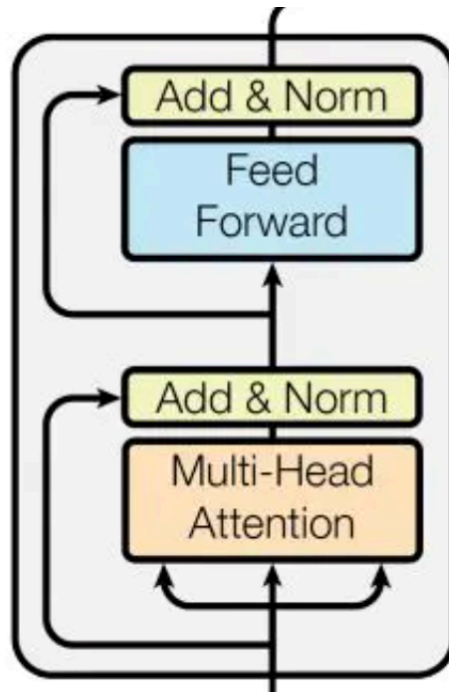
6         self.relu = nn.ReLU()
7         self.linear2 = nn.Linear(ffn_dim, model_dim)
8         self.dropout = nn.Dropout(dropout)
9         self.layer_norm = nn.LayerNorm(model_dim)
10
11     def forward(self, x):
12         residual = x
13
14         x = self.relu(self.linear1(x))
15         x = self.dropout(self.linear2(x))
16         output = self.layer_norm(x + residual)
17
18     return output

```

Encoder

EncoderLayer

EncoderLayer 的每一层都由两部分组成：第一部分是 Multi-Head Self-Attention Mechanism，第二部分是 Position-Wise Feed Forward Network，这两部分都有一个 Residual Connection 和 Layer Normalization.

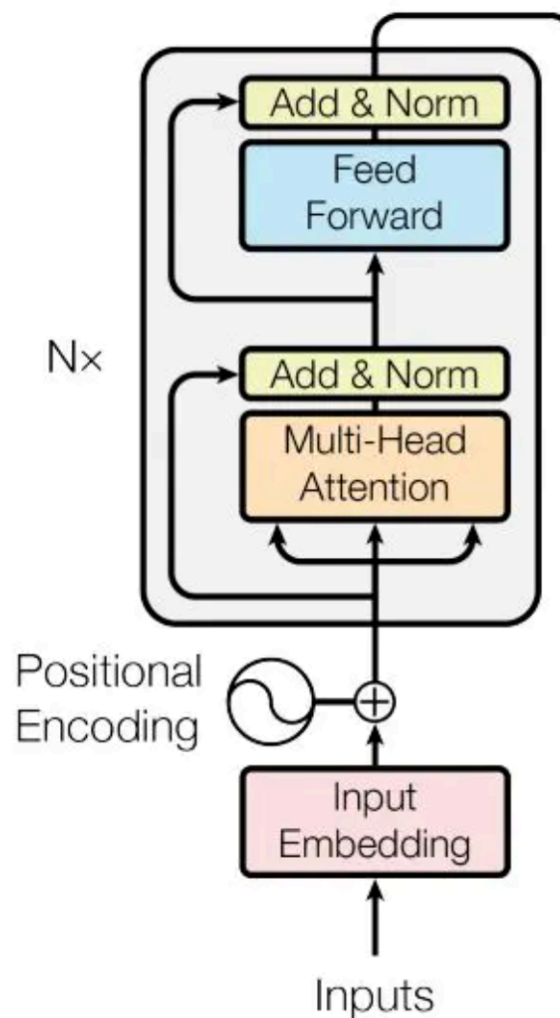



```

1  class EncoderLayer(nn.Module):
2      def __init__(self, model_dim=512, num_heads=8, ffn_dim=2048,
3          dropout=0.2):
4          super(EncoderLayer, self).__init__()
5
6          self.attention = MultiHeadAttention(model_dim, num_heads, dropout)
7          self.feed_forward = PositionalWiseFeedForward(model_dim, ffn_dim,
8              dropout)
9
10         def forward(self, enc_inputs, enc_self_attn_mask=None):
11             context, attention = self.attention(enc_inputs, enc_inputs,
12                 enc_inputs, enc_self_attn_mask)
13             output = self.feed_forward(context)
14
15         return output, attention

```

Encoder



```

1  class Encoder(nn.Module):
2      def __init__(self, vocab_size, max_seq_length, num_layers=6,
3          model_dim=512,

```

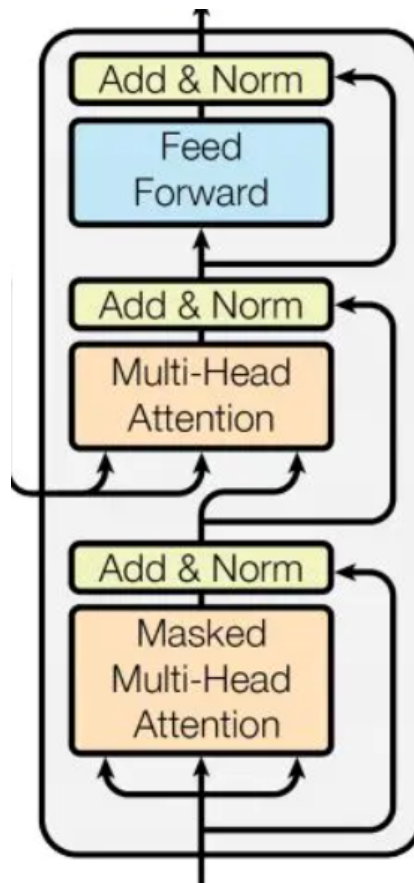
```

3         num_heads=8, ffn_dim=2048, dropout=0.2):
4         super(Encoder, self).__init__()
5
6         self.word_embed = nn.Embedding(vocab_size, model_dim,
padding_idx=0)
7         self.pos_embed = PositionalEncoding(max_seq_length, model_dim)
8         self.dropout = nn.Dropout(dropout)
9         self.layer_norm = nn.LayerNorm(model_dim)
10        self.layer_stack = nn.ModuleList([EncoderLayer(model_dim,
num_heads, ffn_dim, dropout)
11                                           for _ in range(num_layers)])
12
13    def forward(self, inputs, slf_attn_mask):
14        embed = self.word_embed(inputs) + self.pos_embed(inputs)
15        output = self.layer_norm(self.dropout(embed))
16
17        attentions = []
18        for encoder in self.layer_stack:
19            output, attention = encoder(output, slf_attn_mask)
20            attentions.append(attention)
21
22        return output, attentions

```

Decoder

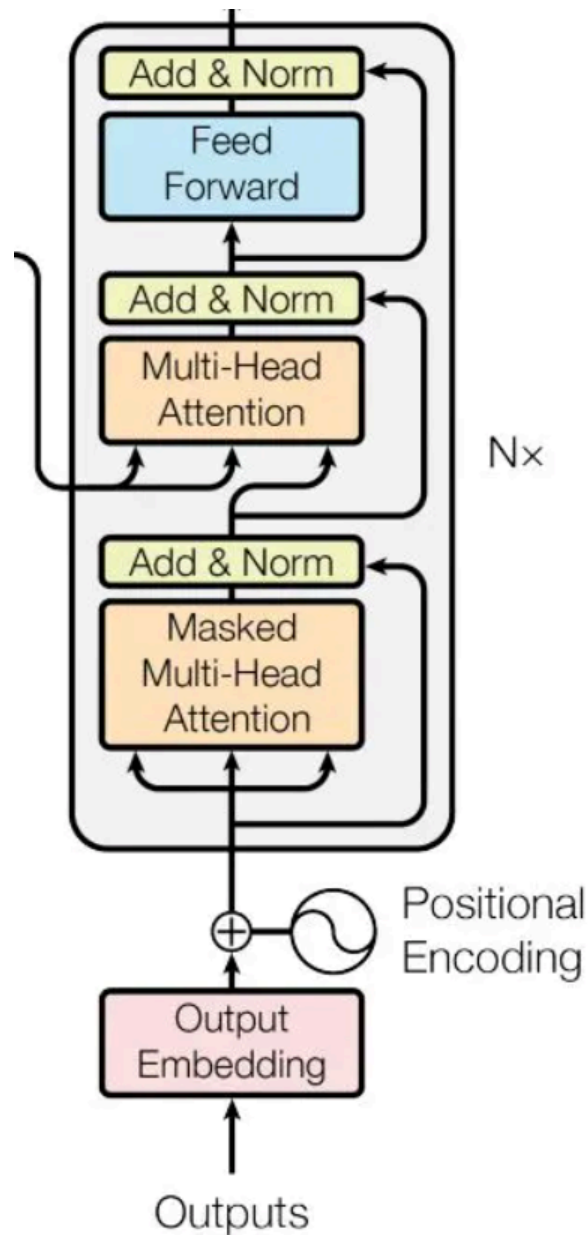
DecoderLayer



DecoderLayer 每一个层包括以下3个部分：第一个部分是 multi-head self-attention mechanism，第二部分是 multi-head context-attention mechanism，第三部分是 position-wise feed-forward network，这三个部分的每一个部分都有一个 residual connection 和 Layer Normalization.

```
1 class DecoderLayer(nn.Module):
2     def __init__(self, model_dim=512, num_heads=8, ffn_dim=2048,
3 dropout=0.2):
4         super(DecoderLayer, self).__init__()
5
6         self.self_attn = MultiHeadAttention(model_dim, num_heads, dropout)
7         self.dec_enc_attn = MultiHeadAttention(model_dim, num_heads,
8 dropout)
9         self.feed_forward = PositionalWiseFeedForward(model_dim, ffn_dim,
10 dropout)
11
12     def forward(self, dec_inputs, enc_outputs, slf_attn_mask=None,
13 dec_enc_attn_mask=None):
14         # self attention: all inputs are decoder inputs
15         dec_outputs, self_attn = self.self_attn(dec_inputs, dec_inputs,
16 dec_inputs, slf_attn_mask)
17
18         # context attention: query is decoder's outputs, key and value are
19 encoder's inputs
20         dec_outputs, dec_enc_attn = self.dec_enc_attn(dec_outputs,
21 enc_outputs, enc_outputs, dec_enc_attn_mask)
22
23         dec_outputs = self.feed_forward(dec_outputs)
24
25         return dec_outputs, self_attn, dec_enc_attn
```

Decoder



```

1  class Decoder(nn.Module):
2      def __init__(self, vocab_size, max_seq_length, num_layers=6,
3          model_dim=512,
4          num_heads=8, ffn_dim=2048, dropout=0.2):
5          super(Decoder, self).__init__()
6
7          self.word_embed = nn.Embedding(vocab_size, model_dim,
8              padding_idx=0)
9          self.pos_embed = PositionalEncoding(max_seq_length, model_dim)
10         self.dropout = nn.Dropout(dropout)
11         self.layer_norm = nn.LayerNorm(model_dim)
12         self.layer_stack = nn.ModuleList([DecoderLayer(model_dim,
13             num_heads, ffn_dim, dropout)
14             for _ in range(num_layers)])
15
16     def forward(self, tgt_seq, tgt_mask, enc_outputs, src_mask):
17         embed = self.word_embed(tgt_seq) + self.pos_embed(tgt_seq)

```

```
15         dec_outputs = self.layer_norm(self.dropout(embed))
16
17         self_attns, dec_enc_attns = [], []
18         for decoder in self.layer_stack:
19             dec_outputs, self_attn, dec_enc_attn = decoder(dec_outputs,
20 enc_outputs, slf_attn_mask=tgt_mask, dec_enc_attn_mask=src_mask)
21             self_attns.append(self_attn)
22             dec_enc_attns.append(dec_enc_attn)
23
24         return dec_outputs, self_attns, dec_enc_attns
```

参考资料

<https://juejin.im/post/5b9f1af0e51d450e425eb32d#heading-13>

<https://github.com/jadore801120/attention-is-all-you-need-pytorch>